

CMSE428 Task 1

November 11, 2020

0.0.1 Çağıl Peköz 17330172

1 Task 1

```
[149]: %matplotlib inline

import pandas as pd
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

1.1 Dataset Loaded in

```
[150]: dataset = pd.read_csv('dataset.csv')
```

1.2 Original Dataset

```
[151]: dataset
```

```
[151]:
```

	ClassLabel	Gender	Age	Race	Education	FamilyHistory	HighBloodSugar	\
0	Negative	1	44	3	3.0	1	1	
1	Negative	2	44	1	2.0	2	1	
2	Negative	1	22	3	5.0	1	1	
3	Negative	2	35	3	4.0	1	1	
4	Negative	1	54	2	5.0	2	1	
...	
1498	Positive	2	59	3	3.0	2	1	
1499	Positive	2	74	3	NaN	2	1	
1500	Positive	2	80	3	2.0	1	1	
1501	Positive	2	78	3	4.0	2	1	
1502	Positive	1	80	2	3.0	1	1	

	BMI	WaistCircumference	SystolicBP	DiastolicBP	LDL	HDL	\
0	29.10		106.6	119	88	113	37
1	25.06		82.3	130	86	138	63
2	28.07		94.5	120	87	96	57
3	20.29		71.9	110	64	93	89

4	36.32	116.4	92	42	110	62
...
1498	26.67	103.0	110	60	132	47
1499	25.08	95.5	135	60	114	36
1500	NaN	119.3	121	40	111	38
1501	46.98	129.5	146	78	136	56
1502	27.14	96.7	154	30	128	65

	Triglycerides
0	235
1	61
2	111
3	58
4	71
...	...
1498	190
1499	270
1500	131
1501	144
1502	108

[1503 rows x 14 columns]

1.3 Checking which columns have missing Data

By running the below code, we can see that Education and BMI has missing data.

```
[152]: dataset.isnull().sum()
```

```
[152]: ClassLabel      0
Gender                0
Age                  0
Race                 0
Education            12
FamilyHistory        0
HighBloodSugar       0
BMI                   9
WaistCircumference   0
SystolicBP           0
DiastolicBP          0
LDL                  0
HDL                  0
Triglycerides        0
dtype: int64
```

1.4 For example, 150th row of BMI has no data

```
[153]: dataset.loc[150, 'BMI']
```

```
[153]: nan
```

1.5 Filling in missing BMI data by Mean and Education by Most Frequent Value

BMI value can be a float, so we can use mean to impute missing data. By running the print command, we can see that the missing values are filled in with mean.

```
[154]: dataset['BMI'].fillna(dataset['BMI'].mean(), inplace=True)
dataset['Education'].fillna(dataset['Education'].value_counts().index[0],
                             ↪inplace=True)
dataset.isnull().sum()
```

```
[154]: ClassLabel      0
Gender              0
Age                0
Race               0
Education          0
FamilyHistory      0
HighBloodSugar     0
BMI                0
WaistCircumference 0
SystolicBP         0
DiastolicBP        0
LDL                0
HDL                0
Triglycerides      0
dtype: int64
```

1.6 Checking the 150th row again

```
[155]: dataset.loc[150, 'BMI']
```

```
[155]: 28.41636546184739
```

1.7 Number of Positives and Negatives

```
[156]: dataset['ClassLabel'].value_counts()
```

```
[156]: Negative      775
Positive      728
Name: ClassLabel, dtype: int64
```

1.8 Number of Positives and Gender==1 and Race==2

```
[157]: len(dataset[(dataset['ClassLabel']=='Positive') & (dataset['Gender']==1) &
↳ (dataset['Race']==2)])
```

[157]: 40

1.9 Race that is most frequent in Negative samples

```
[158]: negativeData = dataset[(dataset['ClassLabel']=='Negative')]
negativeData
```

```
[158]:
```

	ClassLabel	Gender	Age	Race	Education	FamilyHistory	HighBloodSugar	\
0	Negative	1	44	3	3.0	1	1	
1	Negative	2	44	1	2.0	2	1	
2	Negative	1	22	3	5.0	1	1	
3	Negative	2	35	3	4.0	1	1	
4	Negative	1	54	2	5.0	2	1	
..	
770	Negative	2	53	2	5.0	2	1	
771	Negative	1	56	3	5.0	1	1	
772	Negative	2	67	1	4.0	2	1	
773	Negative	2	34	2	1.0	1	1	
774	Negative	2	48	3	5.0	1	1	

	BMI	WaistCircumference	SystolicBP	DiastolicBP	LDL	HDL	\
0	29.10		106.6	119	88	113	37
1	25.06		82.3	130	86	138	63
2	28.07		94.5	120	87	96	57
3	20.29		71.9	110	64	93	89
4	36.32		116.4	92	42	110	62
..
770	20.09		75.0	94	60	148	67
771	30.89		113.1	120	83	139	65
772	20.37		71.6	150	80	95	60
773	38.74		117.4	107	67	91	34
774	21.66		71.5	102	61	88	72

	Triglycerides
0	235
1	61
2	111
3	58
4	71
..	...
770	72
771	84
772	138

```
773          60
774          50
```

```
[775 rows x 14 columns]
```

```
[159]: negativeData['Race'].value_counts()
```

```
[159]: 3    404
      1    127
      4    117
      2     86
      5     41
      Name: Race, dtype: int64
```

1.10 Average BMI of samples with Education==3

```
[160]: education_ft = dataset[(dataset['Education']==3)]
      education_ft['BMI'].mean()
```

```
[160]: 29.276166255235008
```

1.11 Average BMI of Gender==1

```
[161]: gender_filter = dataset[(dataset['Gender']==1)]
      gender_filter['BMI'].mean()
```

```
[161]: 28.49773850117451
```

1.12 Average BMI of Gender==2

```
[162]: gender_filter = dataset[(dataset['Gender']==2)]
      gender_filter['BMI'].mean()
```

```
[162]: 28.347488282367802
```

2 Task 2

2.1 Seperating Categorical and Numerical

Because that datatypes of columns are given not correctly, the columns are seperated manually.

Seperating Categorical Columns

```
[163]: categorical_dataset =
      ↪dataset[['ClassLabel', 'Gender', 'Race', 'Education', 'FamilyHistory', 'HighBloodSugar']]
      categorical_dataset
```

```
[163]:      ClassLabel  Gender  Race  Education  FamilyHistory  HighBloodSugar
0      Negative      1     3      3.0          1            1
1      Negative      2     1      2.0          2            1
2      Negative      1     3      5.0          1            1
3      Negative      2     3      4.0          1            1
4      Negative      1     2      5.0          2            1
...      ...      ...      ...      ...      ...      ...
1498    Positive      2     3      3.0          2            1
1499    Positive      2     3      4.0          2            1
1500    Positive      2     3      2.0          1            1
1501    Positive      2     3      4.0          2            1
1502    Positive      1     2      3.0          1            1
```

[1503 rows x 6 columns]

Seperating Numerical Values

```
[164]: numerical_dataset = \
    ↳ dataset[['Age', 'BMI', 'WaistCircumference', 'SystolicBP', 'DiastolicBP', 'LDL', 'HDL', 'Triglycer
numerical_dataset
```

```
[164]:      Age      BMI  WaistCircumference  SystolicBP  DiastolicBP  LDL  HDL  \
0      44  29.100000          106.6          119          88  113  37
1      44  25.060000          82.3          130          86  138  63
2      22  28.070000          94.5          120          87   96  57
3      35  20.290000          71.9          110          64   93  89
4      54  36.320000          116.4          92          42  110  62
...      ...      ...      ...      ...      ...      ...
1498    59  26.670000          103.0          110          60  132  47
1499    74  25.080000          95.5          135          60  114  36
1500    80  28.416365          119.3          121          40  111  38
1501    78  46.980000          129.5          146          78  136  56
1502    80  27.140000          96.7          154          30  128  65
```

```
      Triglycerides
0          235
1           61
2          111
3           58
4           71
...      ...
1498          190
1499          270
1500          131
1501          144
1502          108
```


[1503 rows x 8 columns]

2.2 Boxplots for Classes with respect to each Variable

Age


```
[165]: sns.set_style("darkgrid")  
sns.boxplot(data=dataset, x='ClassLabel', y=numerical_dataset.columns[0])
```

```
[165]: <AxesSubplot:xlabel='ClassLabel', ylabel='Age'>
```




```
[166]: sns.boxplot(data=dataset, x='ClassLabel', y=numerical_dataset.columns[1])
```

```
[166]: <AxesSubplot:xlabel='ClassLabel', ylabel='BMI'>
```




```
[167]: sns.boxplot(data=dataset, x='ClassLabel', y=dataset.columns[2])
```

```
[167]: <AxesSubplot:xlabel='ClassLabel', ylabel='WaistCircumference'>
```



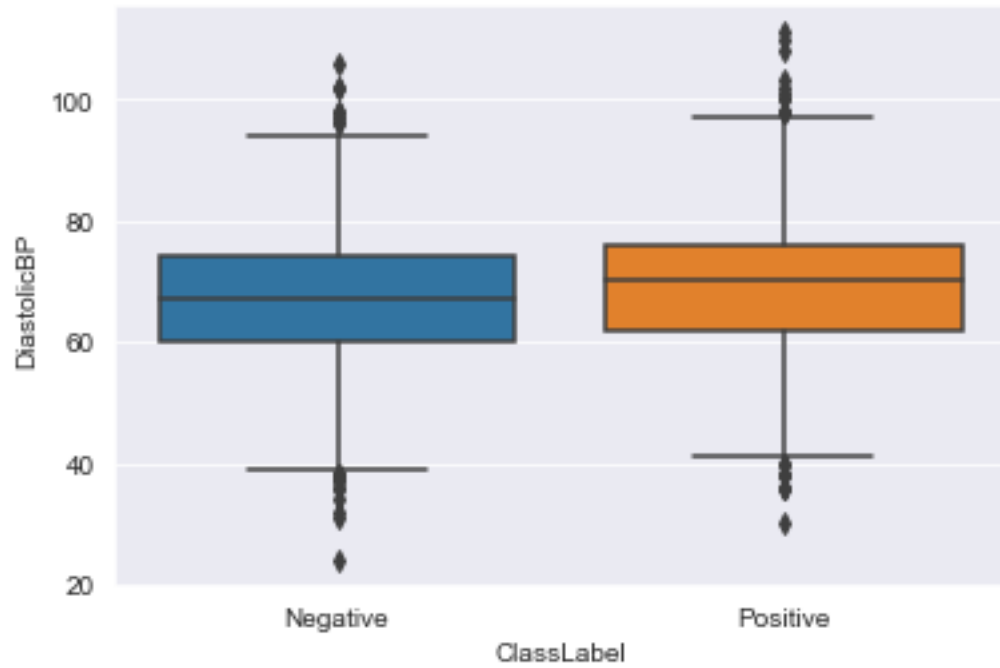

```
[168]: sns.boxplot(data=dataset, x='ClassLabel', y=numerical_dataset.columns[3])
```

```
[168]: <AxesSubplot:xlabel='ClassLabel', ylabel='SystolicBP'>
```



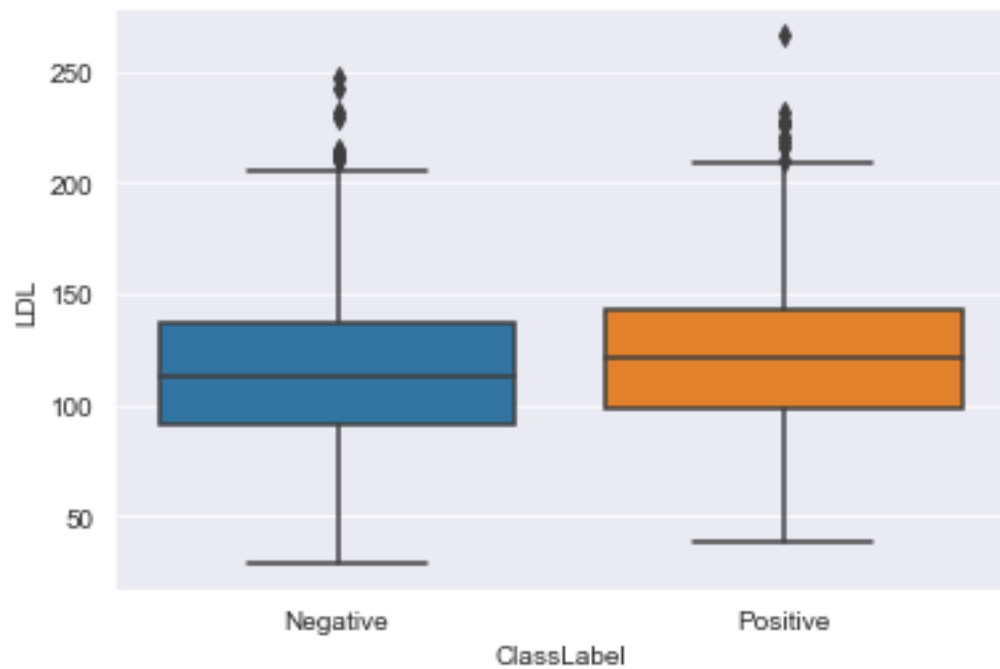
```
[169]: sns.boxplot(data=dataset, x='ClassLabel', y=numerical_dataset.columns[4])
```

```
[169]: <AxesSubplot:xlabel='ClassLabel', ylabel='DiastolicBP'>
```




```
[170]: sns.boxplot(data=dataset, x='ClassLabel', y=numerical_dataset.columns[5])
```

```
[170]: <AxesSubplot:xlabel='ClassLabel', ylabel='LDL'>
```




```
[171]: sns.boxplot(data=dataset, x='ClassLabel', y=numerical_dataset.columns[6])
```

```
[171]: <AxesSubplot:xlabel='ClassLabel', ylabel='HDL'>
```



```
[172]: sns.boxplot(data=dataset, x='ClassLabel', y=numerical_dataset.columns[7])
```

```
[172]: <AxesSubplot:xlabel='ClassLabel', ylabel='Triglycerides'>
```



2.3 Table of Quartiles (Q1,Q2/Median,Q3) for each numerical value

```
[173]: q1_list = [] ; q2_list = [] ; q3_list = []

for column in numerical_dataset:
    q1 = dataset[column].quantile(0.25)
    q2 = dataset[column].mean()
    q3 = dataset[column].quantile(0.75)
    q1_list.append(q1)
    q2_list.append(q2)
    q3_list.append(q3)

quartiles = {}
quartiles['index']=numerical_dataset.columns.tolist()
quartiles['Q1']=q1_list
quartiles['Q2']=q2_list
quartiles['Q3']=q3_list

quartiles_df = pd.DataFrame(quartiles, index=quartiles['index'],
    ↳columns=['Q1', 'Q2', 'Q3'])
quartiles_df
```

```
[173]:
```

	Q1	Q2	Q3
Age	34.00	49.231537	63.000

BMI	24.16	28.416365	31.715
WaistCircumference	86.50	97.555888	106.350
SystolicBP	108.00	120.359947	130.000
DiastolicBP	61.00	68.152362	76.000
LDL	94.00	118.815702	140.000
HDL	44.00	55.480373	65.000
Triglycerides	75.00	117.979375	146.000

2.4 Number of Outliers for each continous variable

```
[174]: print('---Amount of Outliers---')
for index in quartiles_df.index:
    q1 = quartiles_df.loc[index, 'Q1']
    q2 = quartiles_df.loc[index, 'Q2']
    q3 = quartiles_df.loc[index, 'Q3']

    iqr = q3 - q1

    upper_limit = q3 + 1.5*iqr
    lower_limit = q1 - 1.5*iqr

    print (index+' : '+str(len(dataset[(dataset[index]> upper_limit) |
    ↪(dataset[index] < lower_limit)])))
```


```
---Amount of Outliers---
Age: 0
BMI: 42
WaistCircumference: 22
SystolicBP: 32
DiastolicBP: 30
LDL: 19
HDL: 24
Triglycerides: 59
```

3 Task 3

3.1 BarPlot for each categorical variable


```
[175]: sns.barplot(x=categorical_dataset.iloc[:, 0].value_counts().index,
    ↪y=categorical_dataset.iloc[:, 0].value_counts())
```

```
[175]: <AxesSubplot:ylabel='ClassLabel'>
```




```
[176]: sns.barplot(x=categorical_dataset.iloc[:, 1].value_counts().index,   
                  ↪y=categorical_dataset.iloc[:, 1].value_counts())
```

```
[176]: <AxesSubplot:ylabel='Gender'>
```




```
[177]: sns.barplot(x=categorical_dataset.iloc[:, 2].value_counts().index,   
→y=categorical_dataset.iloc[:, 2].value_counts())
```

```
[177]: <AxesSubplot:ylabel='Race'>
```




```
[178]: sns.barplot(x=categorical_dataset.iloc[:, 3].value_counts().index,   
→y=categorical_dataset.iloc[:, 3].value_counts())
```

```
[178]: <AxesSubplot:ylabel='Education'>
```




```
[179]: sns.barplot(x=categorical_dataset.iloc[:, 4].value_counts().index,
↳ y=categorical_dataset.iloc[:, 4].value_counts())
```

```
[179]: <AxesSubplot:ylabel='FamilyHistory'>
```




```
[180]: sns.barplot(x=categorical_dataset.iloc[:, 5].value_counts().index,   
→y=categorical_dataset.iloc[:, 5].value_counts())
```

```
[180]: <AxesSubplot:ylabel='HighBloodSugar'>
```



4 Task 4

```
[181]: sns.pairplot(numerical_dataset)
```

```
[181]: <seaborn.axisgrid.PairGrid at 0x16891c8bbb0>
```

